

Time Series Fundamentals

A Comprehensive Professional Guide to Time-Based Data Analysis

This guide presents the foundational concepts of Time Series analysis in a clear, structured manner — covering data types, key patterns, stationarity, autocorrelation functions, and a complete range of statistical, machine learning, and deep learning models used for forecasting and prediction.

1. What is Time Series?

A **Time Series** is an ordered sequence of data points collected at successive points in time, typically at uniform intervals. It is one of the most fundamental data structures in statistics, economics, engineering, and data science.

1.1 Defining Characteristics

- ◆ The **order of data matters** — each value is tied to its position in time.
- ◆ Each value **depends on time**; the sequence cannot be shuffled without losing meaning.

1.2 Real-World Examples

- ◆ Daily temperature recordings at a meteorological station
- ◆ Stock price movements recorded every minute on an exchange
- ◆ Industrial IoT sensor readings from manufacturing equipment

2. What is Time Series Data?

Time Series Data is a structured dataset where observations are recorded at **regular, consistent time intervals**. The temporal dimension is the single most important feature of this data type.

2.1 Common Measurement Intervals

- ◆ **Hourly** — energy consumption, web server traffic
- ◆ **Daily** — weather readings, website visits, sales figures
- ◆ **Monthly** — revenue reports, production totals
- ◆ **Yearly** — GDP statistics, population census data

★ **Key Point: Time is the primary and most critical feature in any Time Series dataset.**

3. Types of Data

3.1 Cross-Sectional Data

Data collected from **multiple entities at a single point in time**. There is no temporal dependency between observations.

- ◆ Income levels of 100 individuals surveyed in 2024
- ◆ Examination scores of an entire student cohort on one test day

3.2 Pooled (Panel) Data

A combination of Time Series and Cross-Sectional data — tracking **multiple entities across multiple time periods**.

- ◆ Annual employee salaries tracked over five consecutive years
- ◆ Monthly sales performance across multiple retail store branches

Data Type	Description
Cross-Sectional	Multiple entities observed at one single point in time
Time Series	One entity observed across many successive time points
Pooled (Panel)	Multiple entities tracked across multiple time periods

4. Patterns in Time Series

Identifying the underlying patterns in Time Series data is essential for selecting the right model and producing accurate forecasts. There are four main pattern types:

4.1 Trend

A **long-term directional movement** — either upward or downward — in the data over an extended period. A trend reflects the overall direction of the series, independent of short-term fluctuations.

4.2 Seasonality

A pattern that **repeats at fixed, regular calendar intervals**, driven by external factors such as weather cycles, public holidays, or business quarters.

- ◆ Ice cream sales spike every summer and fall every winter
- ◆ E-commerce revenue surges during festive seasons (Diwali, Christmas, Black Friday)

★ **Key Point:** If a pattern repeats consistently after a fixed time period, it is classified as **Seasonality**.

4.3 Cyclic Pattern

Irregular, longer-term fluctuations that do **not** follow a fixed calendar schedule. Cyclic patterns are typically driven by economic or business cycles spanning multiple years.

4.4 Noise (Irregular Component)

Random, unpredictable variation in the data that cannot be attributed to trend, seasonality, or cyclic patterns. Also referred to as the **residual component**.

5. Seasonality — In Depth

Seasonality arises from recurring external forces that systematically and predictably affect observed values at regular intervals.

5.1 Primary Drivers

- ◆ Weather patterns (temperature, rainfall, daylight hours)
- ◆ Cultural and religious festivals (holidays, harvest seasons)
- ◆ Business and financial cycles (quarterly results, fiscal year-end)

5.2 How to Identify Seasonality

- ◆ Look for **repeated peaks and troughs** at consistent calendar intervals
- ◆ Verify whether the same wave-like pattern reappears every month or every year

Monthly Sales Data — Seasonality Example

Month	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct	Nov	Dec
Sales	100	120	150	200	250	300	320	300	250	200	180	150

★ **Key Point:** Sales rise through summer and decline through winter — and this pattern repeats every year, confirming Seasonality.

6. Stationary vs Non-Stationary Data

A Time Series is considered **stationary** if its statistical properties remain constant over time. Stationarity is a foundational requirement for many classical forecasting models.

6.1 Properties That Must Remain Constant

- ◆ **Mean (μ)** — the average value does not drift upward or downward
- ◆ **Variance (σ^2)** — the spread of values remains consistent
- ◆ **Autocorrelation** — the relationship between values at different lags is stable

6.2 Non-Stationary Indicators

- ◆ A visible upward or downward **trend** in the series
- ◆ **Changing variance** — values spread out more or less over time
- ◆ Clear **seasonality** present in the raw data

Type	Example	Indicator
Non-Stationary	100 → 120 → 150 → 200 → 260	Increasing trend — mean is changing
Stationary	100 → 105 → 98 → 102 → 101	Fluctuates around a constant mean

★ **Key Point:** Most classical models such as ARIMA require stationary data to produce reliable forecasts.

6.3 How to Test Stationarity

Method 1 — Visual Inspection

Plot the time series. A visible trend or repeating seasonal shape confirms non-stationarity.

Method 2 — Augmented Dickey-Fuller (ADF) Test

A formal statistical hypothesis test for stationarity:

- ◆ Null Hypothesis: The series is non-stationary
- ◆ If **p-value < 0.05** → reject the null → the data is stationary

Method 3 — Rolling Statistics

Calculate the rolling mean and rolling variance over a sliding window. If both remain approximately constant, the series is stationary.

7. ACF & PACF

7.1 Autocorrelation Function (ACF)

ACF quantifies the **total correlation** between a Time Series and its lagged versions, capturing both direct and indirect relationships across time.

- ◆ Example: How strongly is today's value correlated with yesterday's, two days ago, etc.?

7.2 Partial Autocorrelation Function (PACF)

PACF measures the **direct relationship** between the current value and a specific lag, after statistically removing the influence of all intermediate lags.

- ◆ Lag = the number of time steps back being examined

7.3 Key Differences

Function	Measures	Used To Identify
ACF	Total correlation (direct + indirect)	MA order (q)
PACF	Direct correlation only	AR order (p)

8. Time Series Models

Time Series models are mathematical frameworks trained on historical data to identify patterns and generate reliable predictions of future values.

8.1 Statistical Models

AR — Autoregressive Model

Predicts future values using a **weighted linear combination of past observed values**. The model assumes the current value is a function of its own previous values.

$$Y(t) = c + a_1 \cdot Y(t-1) + a_2 \cdot Y(t-2) + \dots + \text{error}$$

- ◆ c → constant/intercept term
- ◆ a_1, a_2 → learned weights for each lag
- ◆ $Y(t-1), Y(t-2)$ → past observations (Lag 1, Lag 2)

★ **Key Point: AR uses past values only. It does not account for past prediction errors.**

MA — Moving Average Model

Predicts future values using **past forecast errors** rather than past observed values. It smooths the series by correcting for previous mistakes.

$$Y(t) = \text{mean} + b_1 \cdot \text{error}(t-1) + b_2 \cdot \text{error}(t-2) + \dots$$

★ **Key Point: MA uses past errors only — not past values.**

ARMA Model

A hybrid model combining AR (past values) and MA (past errors) components. Effective and well-suited for **stationary data only**.

ARIMA Model

ARIMA extends ARMA to handle **non-stationary data** by introducing a differencing step (the 'I' component) that transforms the series into a stationary one before modelling.

$$\text{ARIMA} = \text{AR (past values)} + \text{I (differencing)} + \text{MA (past errors)}$$

- ◆ Step 1: Detect non-stationarity in the raw data
- ◆ Step 2: Apply differencing to achieve stationarity
- ◆ Step 3: Fit an ARMA model on the now-stationary series

★ **Key Point: ARIMA is the standard choice for real-world data with trend but without seasonality.**

ARIMAX Model

ARIMAX enriches ARIMA by incorporating **external (exogenous) predictor variables** that influence the target series beyond its own history.

$$\text{ARIMAX} = \text{ARIMA} + \text{External Variables (X)}$$

- ◆ Example: Predicting sales using past sales + advertising spend + promotions

SARIMA Model

SARIMA extends ARIMA by adding explicit **seasonal AR, I, and MA terms**, making it capable of modelling repeating seasonal patterns.

$$\text{SARIMA} = \text{ARIMA} + \text{Seasonal Terms (P, D, Q)}_m$$

- ◆ Example: Monthly retail sales exhibiting a strong annual seasonal cycle

SARIMAX Model

The most comprehensive statistical framework — combining seasonality, differencing, AR/MA components, and external predictor variables.

SARIMAX = SARIMA + External Variables (X)

◆ Example: Predicting seasonal sales using history + festival data + marketing campaigns

★ **Key Point: SARIMAX handles seasonality AND external drivers — ideal for complex real-world scenarios.**

8.2 Machine Learning Models

XGBoost

A powerful **gradient-boosted tree ensemble** algorithm adapted for Time Series by engineering lag features, rolling statistics, and calendar/date features as inputs.

8.3 Deep Learning Models

RNN — Recurrent Neural Network

A neural network architecture designed for sequential data. Each output step is fed back as input to the next, creating an internal memory of recent history.

LSTM — Long Short-Term Memory

An advanced RNN variant using gating mechanisms (input, forget, output gates) to capture **long-term dependencies**, overcoming the vanishing gradient problem.

GRU — Gated Recurrent Unit

A streamlined version of LSTM with fewer parameters, offering comparable performance with greater computational efficiency in most tasks.

Autoencoder

Learns a **compressed latent representation** of the input series. Particularly effective for unsupervised **anomaly detection** in time series data.

VAE — Variational Autoencoder

A probabilistic autoencoder that models the latent space as a distribution, enabling uncertainty-aware representation learning and generation.

8.4 Modern Frameworks

Prophet

An open-source forecasting library developed by Meta (Facebook). It automatically detects and handles **trend changes, multiple seasonalities, and holiday effects** with minimal configuration — accessible to analysts without deep statistical expertise.

9. Model Comparison — Quick Reference

The table below summarises all Time Series models covered in this guide, along with their category and optimal use cases.

Model	Category	Best Use Case
AR	Statistical	Stationary data with strong autocorrelation
MA	Statistical	Stationary data, correcting for past errors
ARMA	Statistical	Stationary data — combines AR + MA
ARIMA	Statistical	Non-stationary data with trend (no seasonality)
ARIMAX	Statistical	Non-stationary + external predictor variables
SARIMA	Statistical	Seasonal non-stationary time series
SARIMAX	Statistical	Seasonal data + external variables
XGBoost	Machine Learning	Feature-rich tabular time series (lag features)
LSTM / GRU	Deep Learning	Long-range sequential dependencies
Autoencoder	Deep Learning	Anomaly detection in sequences
Prophet	Framework	Auto trend + seasonality, minimal configuration

Created by — Priyanka Deore & Kunal Mahadule